



+11,00.00

# NETWORKS PROGRAMMING

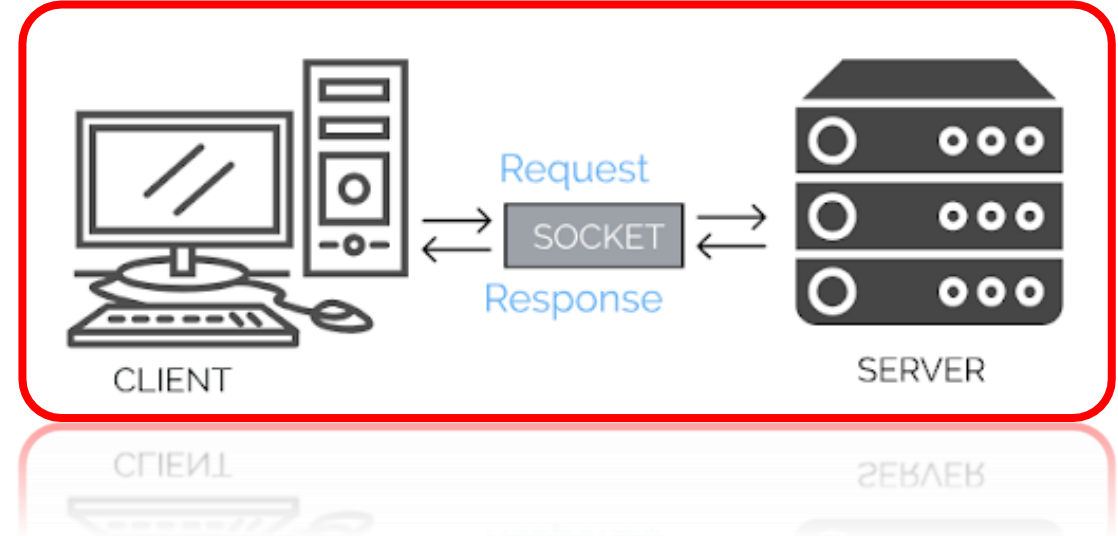
LECTION3 NETWORK API



# LECTION OBJECTIVES

- ❖ What are Sockets in Network Programs
- ❖ IPv4 vs IPv6
- ❖ Byte Order
- ❖ Network API in details
- ❖ Client-Server Communication
- ❖ Practice

A **network socket** is an internal endpoint for sending or receiving data within a node on a computer network.

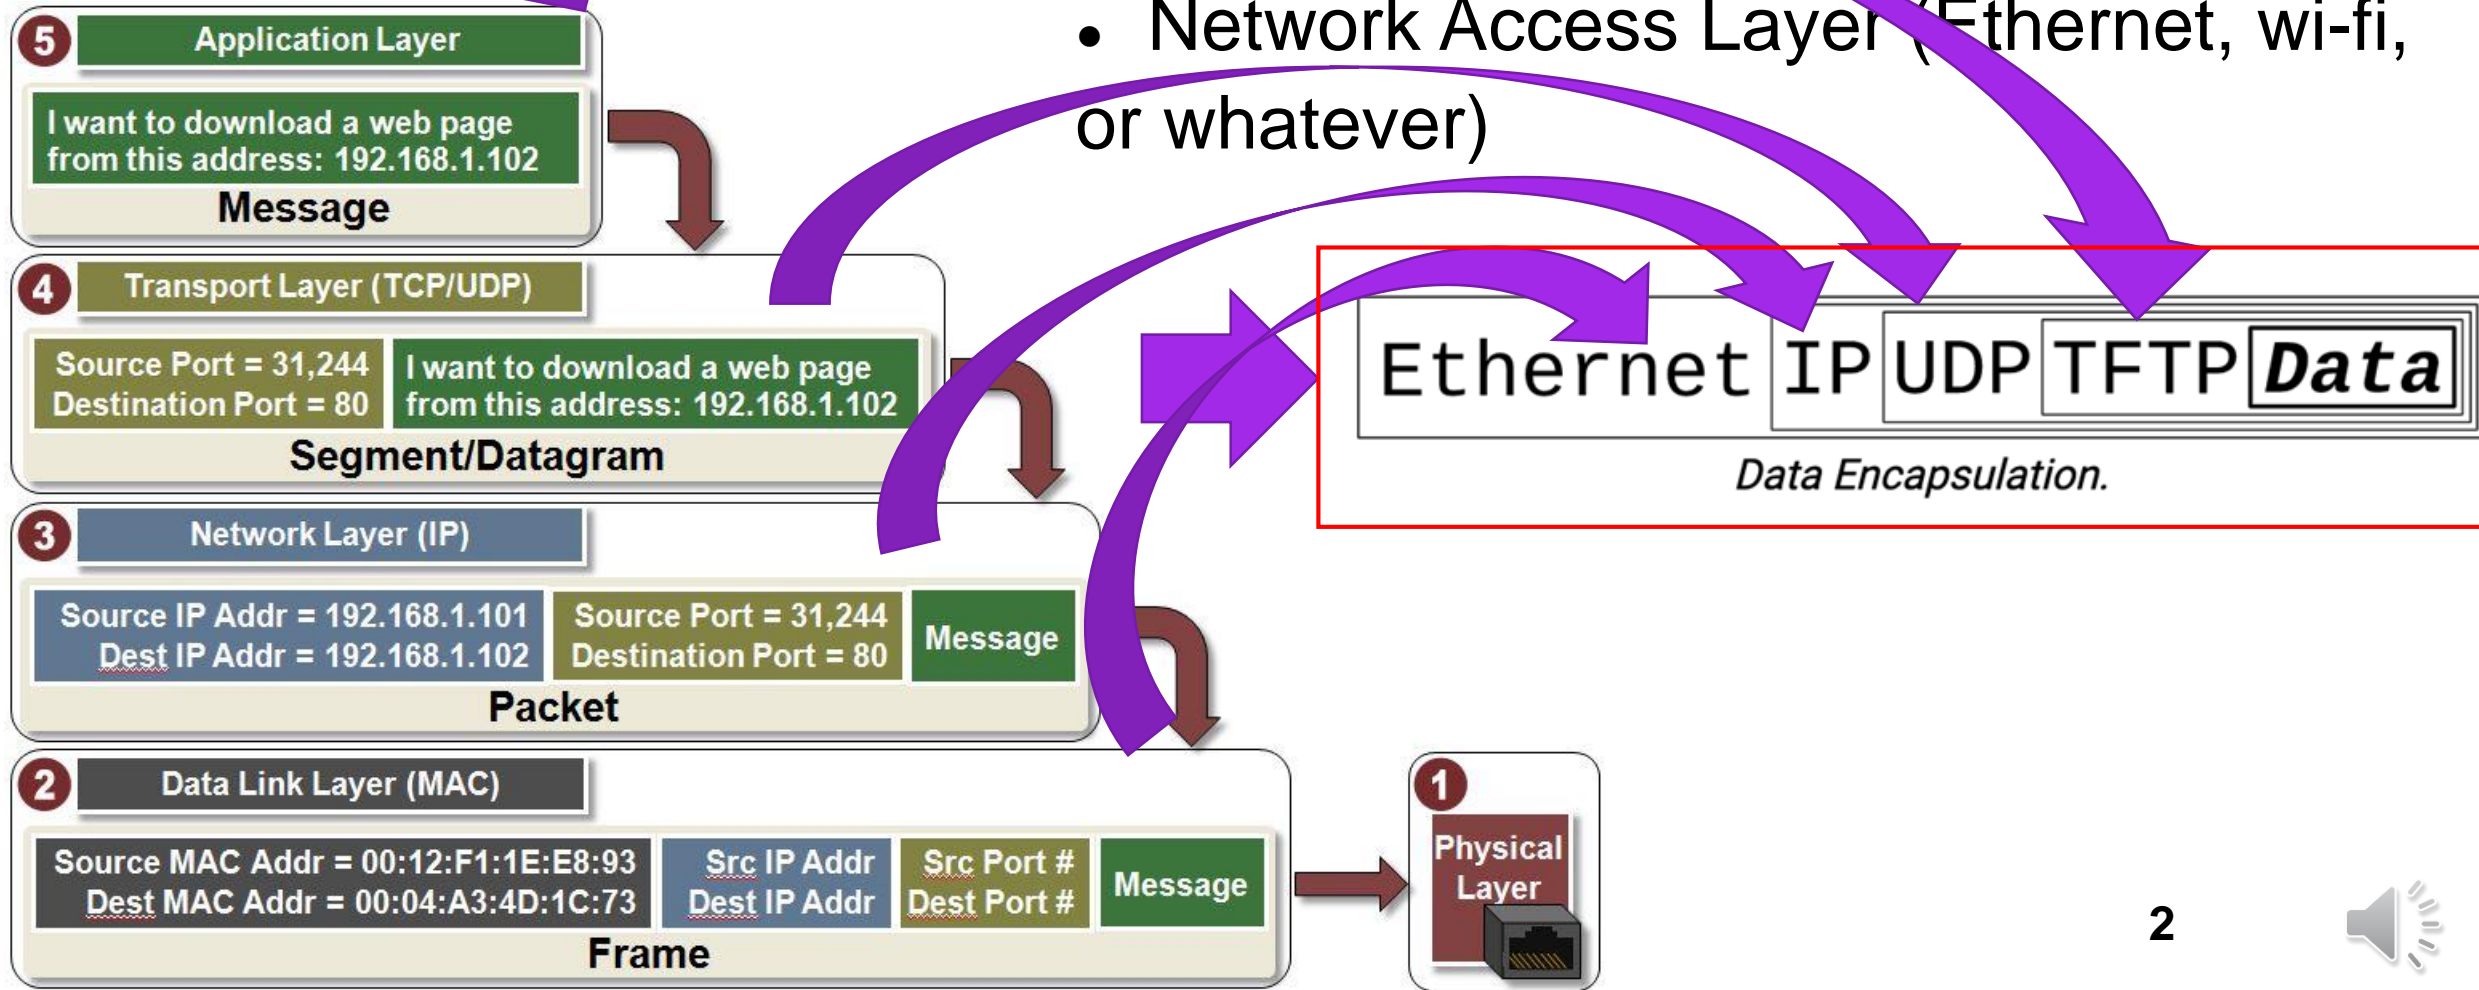


The term *socket* is analogous to physical **female connectors**, communication between two nodes through a **channel** being visualized as a cable with two **male connectors** plugging into sockets at each node.



## Recall Network Theory


- Application Layer (telnet, ftp, etc.)
- Transport Layer (TCP, UDP)
- Internet Layer (IP and routing)
- Network Access Layer (Ethernet, wi-fi, or whatever)



# What is a socket!?

- **Socket:** a way to speak to other programs using standard Unix file descriptors
- **File descriptor** is simply an integer associated with an open file. But (and here's the catch), that file can be a network connection, a FIFO, a pipe, a terminal, a real on-the-disk file, or just about anything else. Everything in Unix is a file!
- Make a call to the **socket()** system routine. It returns the socket descriptor, and you communicate through it using the specialized **send()** and **recv()** (man send, man recv) socket calls.

*socket:* door between application process and end-end-transport protocol



Command for use in the terminal  
get the documentation.

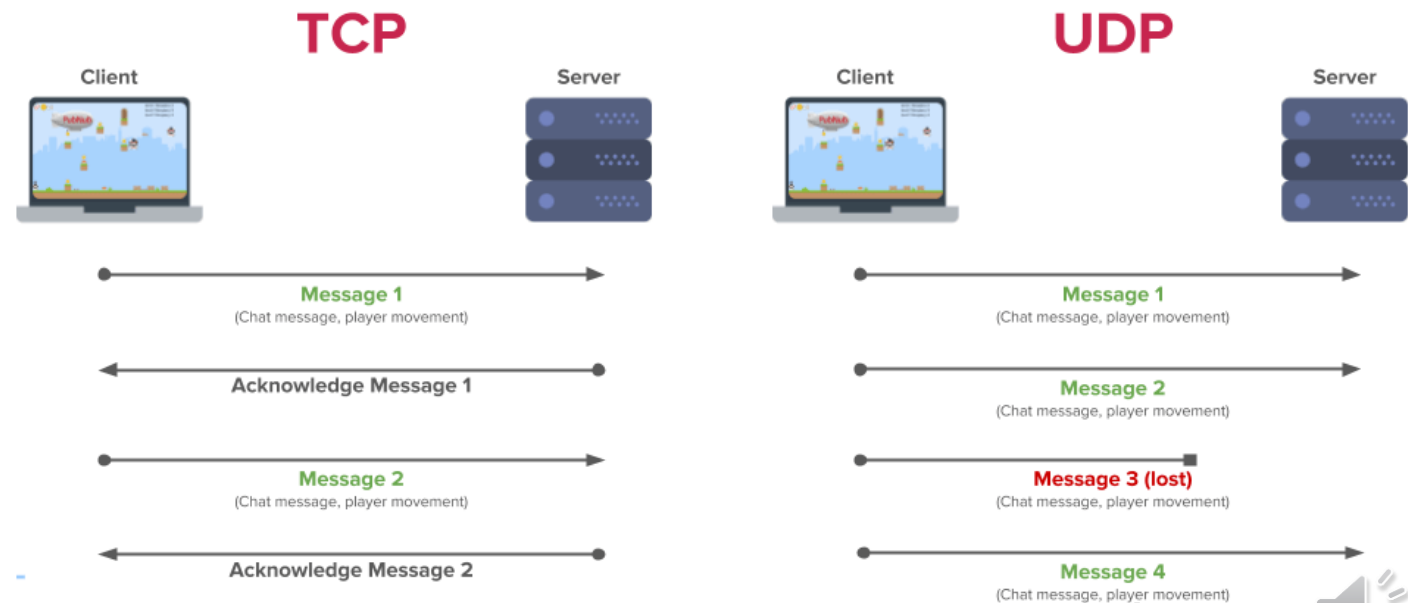





# Two types of internet sockets

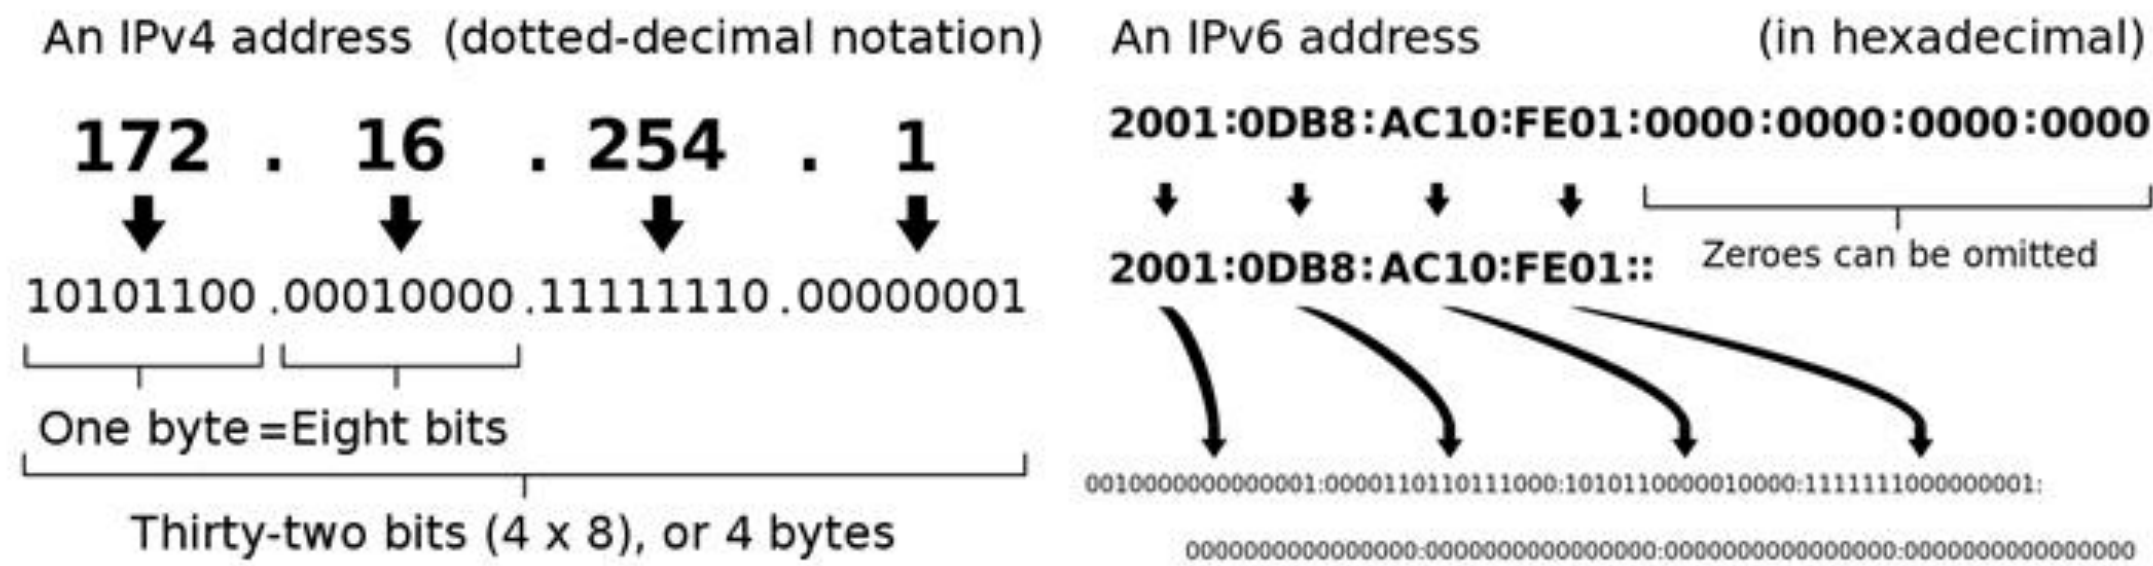
- **SOCK\_STREAM** (TELNET, HTTP, MAIL, ... over TCP) - Provides sequenced, reliable, two-way, connection-based byte streams. An out-of-band data transmission mechanism may be supported
- **SOCK\_DGRAM** (TFTP, DHCP, video streaming, audio streaming, ...) - supports datagrams (connectionless, unreliable messages of a fixed maximum length)

• ...




# IP address v4/v6

- 192.0.2.111 - IPv4
- 2001:0db8:c9d2:aee5:73e3:934a:a5ae:9551 - IPv6
- The address ::1 is the loopback address. It always means “this machine I’m running on now”. In IPv4, the loopback address is 127.0.0.1.




# The IPv6 Address Space

- 128-bit address space
- 128 bits were chosen to allow multiple levels of hierarchy and flexibility in designing hierarchical addressing and routing
- Global unicast and anycast addresses are defined by a global routing prefix, a subnet ID, and an interface ID




# IP address v4/v6

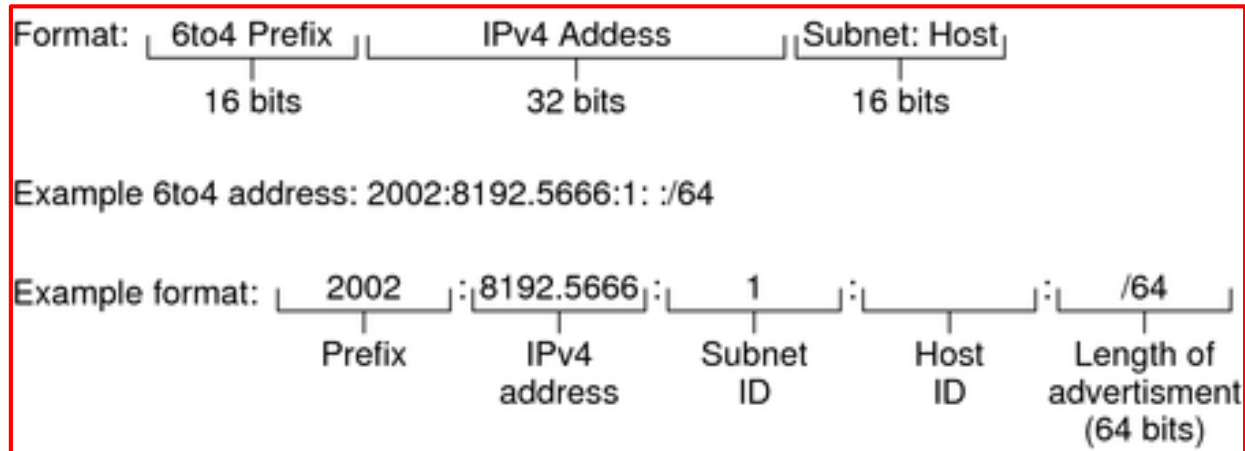
- Subnets - Mask & Class



Before  
subnetting



After  
subnetting


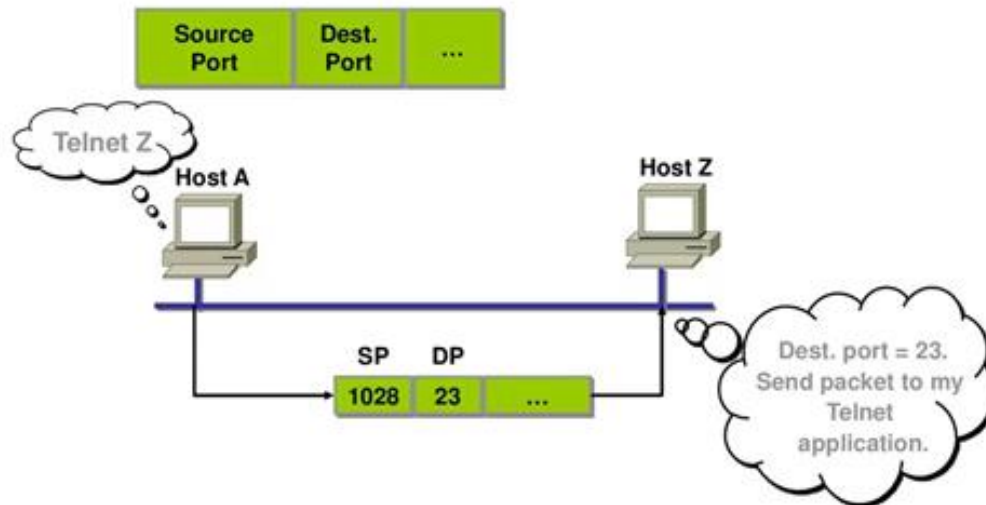




# Port numbers

- TCP/UDP (transport level)
- **Port number** is a 16-bit number that's like the local address for the connection
- HTTP 80, TELNET 23, SMTP 25, DOOM 666
- < 1024 reserved

## TCP Port Numbers



IP Address + Port number = Socket

## TCP/IP Ports And Sockets

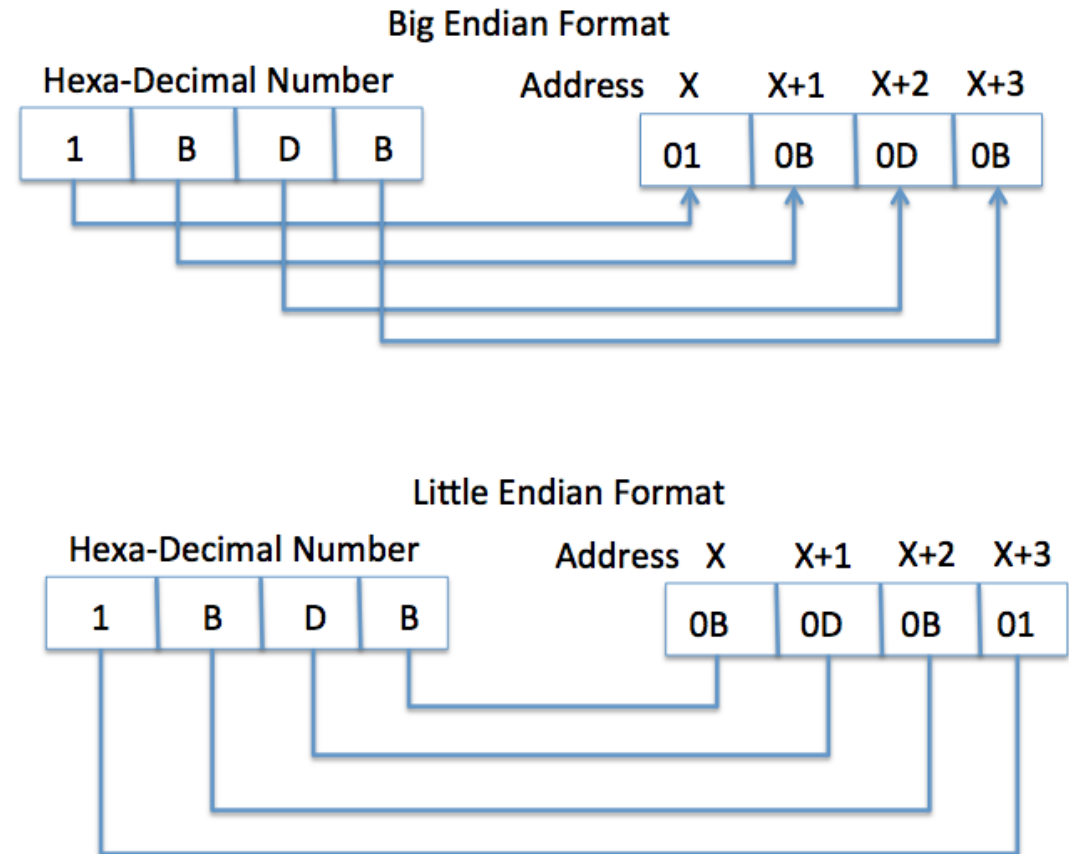


# Byte Order

- Computer might have been storing bytes in reverse order behind your back
- **Big-Endian** is also called Network Byte Order because that's the order us network types like
- **Little-Endian** stores the least-significant byte at the smallest

<u>Function</u>	<u>Description</u>
htons()	host to network short
htonl()	host to network long
ntohs()	network to host short
ntohl()	network to host long

## Big Endian and Little Endian Formats



# Structures

```
struct addrinfo {  
    int          ai_flags;           // AI_PASSIVE, AI_CANONNAME, etc.  
    int          ai_family;         // AF_INET, AF_INET6, AF_UNSPEC  
    int          ai_socktype;       // SOCK_STREAM, SOCK_DGRAM  
    int          ai_protocol;       // use 0 for "any"  
    size_t       ai_addrlen;        // size of ai_addr in bytes  
    struct sockaddr *ai_addr;       // struct sockaddr_in or _in6  
    char         *ai_canonname;     // full canonical hostname  
  
    struct addrinfo *ai_next;       // linked list, next node  
};
```



# Structures

```
struct sockaddr {
    unsigned short    sa_family;    // address family, AF_xxx
    char              sa_data[14];  // 14 bytes of protocol address
};

struct sockaddr_in {
    short int         sin_family;    // Address family, AF_INET
    unsigned short int sin_port;     // Port number
    struct in_addr     sin_addr;     // Internet address
    unsigned char      sin_zero[8];  // Same size as struct sockaddr
};

// Internet address (a structure for historical reasons)
struct in_addr {
    uint32_t s_addr; // that's a 32-bit int (4 bytes)
};
```



# Structures

```
struct sockaddr_in6 {  
    u_int16_t      sin6_family;    // address family, AF_INET6  
    u_int16_t      sin6_port;      // port number, Network Byte Order  
    u_int32_t      sin6_flowinfo;  // IPv6 flow information  
    struct in6_addr sin6_addr;      // IPv6 address  
    u_int32_t      sin6_scope_id;  // Scope ID  
};
```

```
struct in6_addr {  
    unsigned char  s6_addr[16];    // IPv6 address  
};
```





# IP addresses conversion From IP4 to IP6

- No need to figure them out by hand and stuff them in a long with the in operator.
- `inet_pton()`, converts an IP address in numbers-and-dots notation into either a struct `in_addr` or a struct `in6_addr` depending on whether you specify `AF_INET` or `AF_INET6`

- `inet_pton()` reverse transformation

```
struct sockaddr_in sa; // IPv4
struct sockaddr_in6 sa6; // IPv6
```

```
inet_pton(AF_INET, "10.12.110.57", &(sa.sin_addr)); // IPv4
inet_pton(AF_INET6, "2001:db8:63b3:1::3490", &(sa6.sin6_addr)); // IPv6
```

```
char ip4[INET_ADDRSTRLEN]; // space to hold the IPv4 string
struct sockaddr_in sa; // pretend this is loaded with something
```

```
inet_ntop(AF_INET, &(sa.sin_addr), ip4, INET_ADDRSTRLEN);
```



# IP addresses conversion From IP4 to IP6

1. First of all, try to **use getaddrinfo()** to get all the struct sockaddr info, instead of packing the structures by hand.
2. Any place that you find you're hard-coding anything related to the IP version, try to wrap up in a helper function.
3. Change **AF\_INET** to **AF\_INET6**.
4. Change **PF\_INET** to **PF\_INET6**.
5. Change the assignments of **INADDR\_ANY** to **IN6ADDR\_ANY**
6. Instead of struct sockaddr\_in **use struct sockaddr\_in6**
7. Instead of struct in\_addr **use struct in6\_addr**
8. Instead of inet\_aton() or inet\_addr(), **use inet\_pton()**.
9. Instead of inet\_ntoa(), **use inet\_ntop()**.
10. Instead of gethostbyname(), **use** the superior **getaddrinfo()**.
11. Instead of gethostbyaddr(), **use** the superior **getnameinfo()** (although gethostbyaddr() can **still work with IPv6**)
12. INADDR\_BROADCAST no longer works. **Use IPv6 multicast** instead.

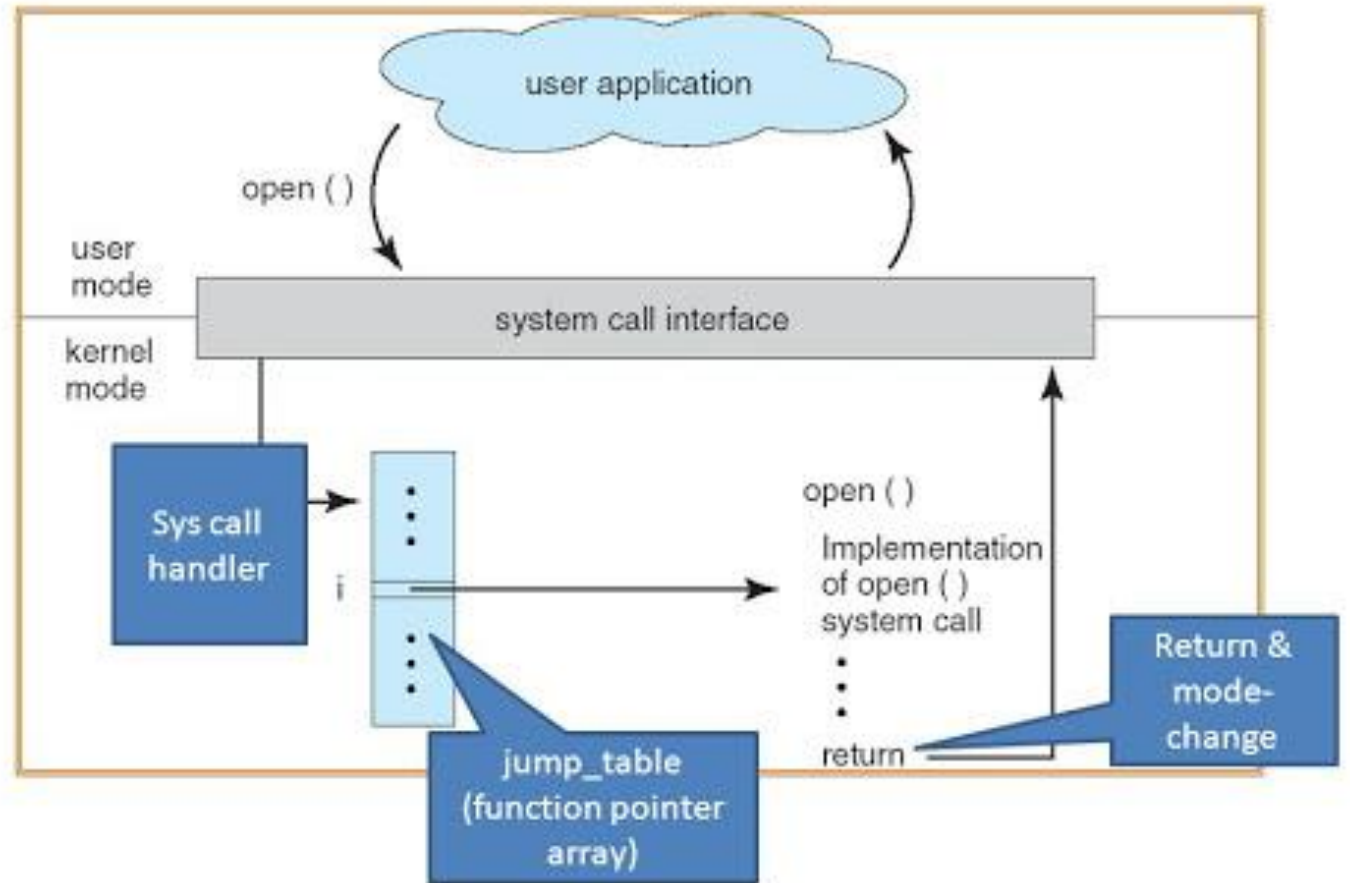


# System calls

**System call** provides the services of the operating system to the user programs via Application Program Interface(API).

It provides an interface between a process and operating system to allow user-level processes to request services of the operating system. System calls are the only entry points into the kernel system

## API – System Call – OS Relationship



# getaddrinfo

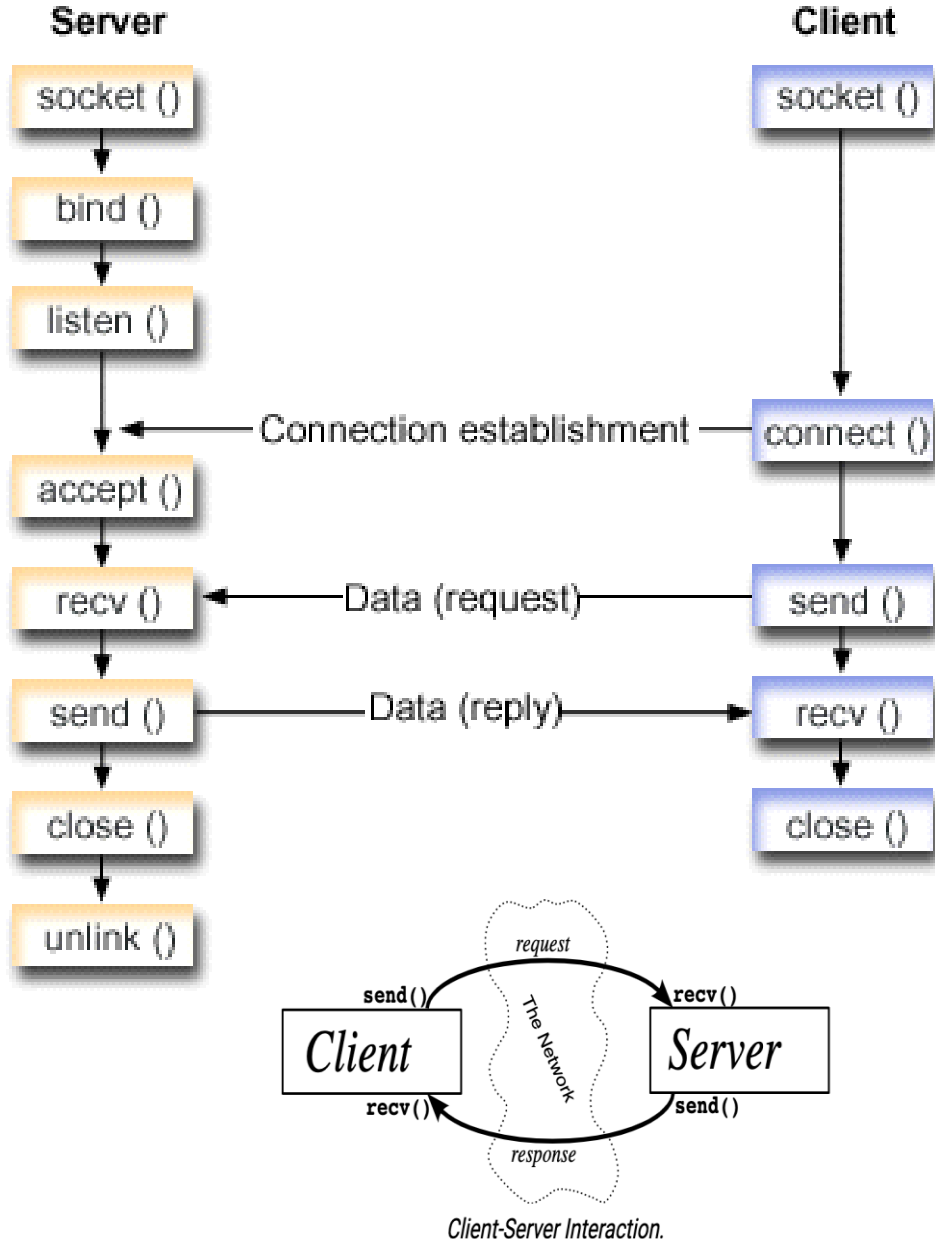
- DNS and service name lookups
- Give this function three input parameters, and it gives you a pointer to a linked-list, `res`, of results.
- The `node` parameter is the host name to connect to, or an IP address.
- `service`, which can be a port number, like “80”, or the name of a particular service like “http” or “ftp” or “telnet” or “smtp” or whatever.
- Finally, the `hints` parameter points to a struct `addrinfo` that you’ve already filled out with relevant information.

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>

int getaddrinfo(const char *node,      // e.g. "www.example.com" or IP
                const char *service,  // e.g. "http" or port number
                const struct addrinfo *hints,
                struct addrinfo **res);
```



## State Diagram for Server and Client Model



**Server program** - a software component of a computing system that performs service (maintenance) functions at the request of a client, giving him access to certain resources or services.

The **client program** interacts with the server using a specific protocol. It can request any data from the server, manipulate data directly on the server, start new processes on the server, etc. The client program can provide data received from the server to the user or use it in some other way, depending on the purpose of the program.

**Socket programming** is a way of connecting two nodes on a network to communicate with each other. One socket (node) listens on a particular port at an IP, while other socket reaches out to the other to form a connection. Server forms the listener socket while client reaches out to the server.





# socket

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol);
```

- Get the File Descriptor
- **domain** is PF\_INET or PF\_INET6
- **type** is SOCK\_STREAM or SOCK\_DGRAM
- **protocol** can be set to 0 to choose the proper protocol for the given type
- **socket()** simply returns to you a socket descriptor that you can use in later system calls, or -1 on error. The global variable **errno** is set to the error's value.



# bind

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int bind(int sockfd, struct sockaddr *my_addr, int addrlen);
```

- Once you have a socket, you might have to associate that socket with a port on your local machine.
- **sockfd** is the socket file descriptor returned by `socket()`. **my\_addr** is a pointer to a `struct sockaddr` that contains information about your address, namely, port and IP address. **addrlen** is the length in bytes of that address.



# connect

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int connect(int sockfd, struct sockaddr *serv_addr, int addrlen);
```

- how to connect to a remote host
- **sockfd** is our friendly neighborhood socket file descriptor
- **serv\_addr** is a struct sockaddr containing the destination port and IP address
- **addrlen** is the length in bytes of the server address structure.



# listen

```
int listen(int sockfd, int backlog);
```

```
getaddrinfo();  
socket();  
bind();  
listen();  
/* accept() goes here */
```

- **sockfd** is the usual socket file descriptor from the `socket()` system call
- **backlog** is the number of connections allowed on the incoming queue
- probably get away with setting it to 5 or 10



# accept

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

- connection will be queued up waiting to be accept()ed
- **sockfd** is the listen()ing socket descriptor
- **addr** will usually be a pointer to a local struct sockaddr\_storage. This is where the information about the incoming connection will go (and with it you can determine which host is calling you from which port)
- **addrlen** is a local integer variable that should be set to sizeof(struct sockaddr\_storage) before its address is passed to accept()





# send/recv

```
int send(int sockfd, const void *msg, int len, int flags);
```

```
int recv(int sockfd, void *buf, int len, int flags);
```

- these two functions are for communicating over stream sockets or connected datagram sockets
- **sockfd** is the socket descriptor you want to send data
- **msg** is a pointer to the data you want to send
- **len** is the length of that data in bytes
- just set flags to 0
- **buf** is the buffer to read the information into
- **len** is the maximum length of the buffer



# sendto/recvfrom

```
int sendto(int sockfd, const void *msg, int len, unsigned int flags,  
           const struct sockaddr *to, socklen_t tolen);
```

```
int recvfrom(int sockfd, void *buf, int len, unsigned int flags,  
             struct sockaddr *from, int *fromlen);
```

- this call is basically the same as the call to send() with the addition of two other pieces of information
- struct sockaddr (which will probably be another struct sockaddr\_in or struct sockaddr\_in6 or struct sockaddr\_storage that you cast at the last minute) which contains the destination IP address and port



# close/shutdown

```
close(sockfd);  
int shutdown(int sockfd, int how);
```

- close the connection on your socket descriptor
- Just in case you want a little more control over how the socket closes, you can use the shutdown
- **sockfd** is the socket file descriptor you want to shutdown, and **how** is one of the following

---

<u>how</u>	<u>Effect</u>
0	Further receives are disallowed
1	Further sends are disallowed
2	Further sends and receives are disallowed (like <code>close()</code> )

---



# gethostname

```
#include <unistd.h>
```

```
int gethostname(char *hostname, size_t size);
```

- It returns the name of the computer that your program is running on
- **hostname** is a pointer to an array of chars that will contain the hostname upon the function's return
- **size** is the length in bytes of the hostname array



# getpeername

```
#include <sys/socket.h>
```

```
int getpeername(int sockfd, struct sockaddr *addr, int *addrlen);
```

- the function **getpeername()** will tell you who is at the other end of a connected stream socket
- **sockfd** is the descriptor of the connected stream socket
- **addr** is a pointer to a struct sockaddr
- **addrlen** is a pointer to an int, that should be initialized to sizeof \*addr or sizeof





# Blocking

```
sockfd = socket(PF_INET, SOCK_STREAM, 0);  
fcntl(sockfd, F_SETFL, O_NONBLOCK);
```

- In a nutshell, “block” is techie jargon for “sleep”
- Lots of functions block: accept, recv
- If you don’t want a socket to be blocking, you have to make a call to fcntl()
- If you put your program in a busy-wait looking for data on the socket, you’ll suck up CPU time like it was going out of style
- A more elegant solution for checking to see if there’s data waiting to be read comes in the following section on poll



# select - synchronous I/O multiplexing

```
#include <sys/time.h>
#include <sys/types.h>
#include <unistd.h>
```

```
int select(int numfds, fd_set *readfds, fd_set *writefds,
           fd_set *exceptfds, struct timeval *timeout);
```

```
struct timeval {
    int tv_sec;      // seconds
    int tv_usec;     // microseconds
};
```

---

<u>Function</u>	<u>Description</u>
FD_SET(int fd, fd_set *set);	Add fd to the set.
FD_CLR(int fd, fd_set *set);	Remove fd from the set.
FD_ISSET(int fd, fd_set *set);	Return true if fd is in the set.
FD_ZERO(fd_set *set);	Clear all entries from the set.

---



# poll - synchronous I/O multiplexing

```
#include <poll.h>
```

```
int poll(struct pollfd fds[], nfds_t nfds, int timeout);
```

```
struct pollfd {  
    int fd;           // the socket descriptor  
    short events;     // bitmap of events we're interested in  
    short revents;    // when poll() returns, bitmap of events that occurred  
};
```

---

<u>Macro</u>	<u>Description</u>
POLLIN	Alert me when data is ready to recv( ) on this socket.
POLLOUT	Alert me when I can send( ) data to this socket without blocking.

---

